

Q1: What is a stream in Snowflake?

A: A stream tracks **change data capture (CDC)** in a table—specifically inserts, updates, and deletes—without altering the source table.

Q2: Can streams be queried directly?

A: Yes, streams behave like tables and can be queried using **SELECT** statements to access change records.

Q3: Do streams store historical data?

A: No. Streams maintain changes only **since the last time** they were consumed. Once consumed, those changes are no longer visible.

Q4: What are the types of streams?

A:

- Standard (default)
- Append-only (only tracks inserts)
- Insert-only (used with external tables)

Q1: What is a task in Snowflake?

A: A **task** automates SQL execution on a **schedule** or based on a **dependency chain**.

Q2: What is the minimum schedule interval for a task?

A: 1 minute.

Q3: Can tasks have dependencies?

A: Yes. Tasks can be linked in **trees** where one task triggers another.

Q4: How can a task consume a stream?

A: A task can reference a stream in its SQL, typically in a **MERGE** or **INSERT** statement, to process captured changes.

Q1: What is the **COPY INTO command used for?**

A: It is used to **load data into** Snowflake tables from files or to **unload data** from tables into external stages.

Q2: What are common sources for **COPY INTO?**

A: Internal stages (e.g., `@my_stage`), Amazon S3, Azure Blob, or Google Cloud Storage.

Q3: What file formats are supported with **COPY INTO?**

A: CSV, JSON, Avro, ORC, Parquet, and XML.

Q4: What clause should you use to skip headers in CSV files?

A: `SKIP_HEADER = 1` in the file format options.

Q1: What kind of architecture does Snowflake use?

A: Multi-cluster shared data architecture with three layers:

- 1 Storage
- 2 Compute (virtual warehouses)
- 3 Cloud services

Q2: What is the role of the cloud services layer?

A: It handles query parsing, optimization, security, metadata, and task orchestration.

Q3: Is data ever duplicated when sharing between accounts?

A: No. Snowflake shares **live, read-only access** without copying data.

Q4: Can compute clusters access each other's cache?

A: No. Each virtual warehouse has **its own cache** and does not share it with others.



WAREHOUSES

Q1: What is a virtual warehouse in Snowflake?

A: A **compute cluster** used to execute queries, load data, and perform DML operations.

Q2: What happens when a warehouse is suspended?

A: It stops consuming credits. It can be resumed automatically or manually.

Q3: What are the warehouse size options?

A: From **X-Small** to **6X-Large**, each doubling the compute power and cost.

Q4: What is auto-scaling?

A: A feature that allows a multi-cluster warehouse to scale out (add clusters) to handle high concurrency.

Q5: Does warehouse size impact query performance?

A: Yes, larger warehouses can process queries faster but consume more credits.